

From selling access to selling work (and what it means for you)

Is it time to say RIP to ARR?



KYLE POYAR

OCT 30, 2024



105



10



7

Share

👋 Hi, it's [Kyle Poyar](#) and welcome to [Growth Unhinged](#), my weekly newsletter exploring the hidden playbooks behind the fastest-growing startups. Subscribe to join **58,177 readers** who get Growth Unhinged delivered to their inbox every Wednesday morning.

Subscribe

Annual recurring revenue (ARR) is the building block of SaaS metrics. And it's the basis of SaaS valuation multiples.

It's (usually) high margin, predictable and growing. Which means SaaS companies are (usually) on track to become highly profitable at scale.

AI throws this off. What's not classic software ARR:

- Charging per successful AI resolution of a support ticket
- Charging per photo edited by AI
- Charging per task completed by the AI agent
- Charging per conversation held by an AI agent
- Charging per input/output token to access an LLM

We're moving away from charging for **access to software** and toward a model of charging for the **work delivered** by a combination of software and AI agents. Let's dive into what's happening and what it means for you.

The rise of disruptive AI pricing models

The interest in next generation pricing models has been truly wild.

Technology companies are realizing they can't solely rely on seat-based subscription in an age of AI, automation and APIs where value is disconnected with how many people are logging in. Perhaps Salesforce going all-in on [Agentforce](#) (and charging per conversation) was the push the industry needed.

I've been tracking some of the most disruptive AI pricing models announced so far; subset are included in the graphic below.

The most pricing



\$2 per conversation
(Agentforce)



INTERCOM

\$0.99 per AI response
(FinAI agent)



INTERCOM

10 free tickets per
per mo. (FinAI copilot)



Per successful auto
resolution (Zendesk)



Microsoft

\$4 per hour of use
(AI copilot for sales)




OpenAI

Per input/output token
(GPT-4o)



Credit-based
(credit = data point or action)



Unlimited AI photo edits



Per task completed by
the AI SDR



Per AI photo retouched



Based on number of
personalized emails



Per demand package
generated by AI



Credit-based (credit =
data points, signals)



Per AI generated summary



Credit-based
(credit = AI workflow)



Per user and editable file
translation



Based on workflow steps
and AI credits



25% per successful
chargeback recovered



Discover more from Kyle Poyar's Growth Unhinged

Revealing the playbooks and hidden tactics behind today's best startups. Join 65,000+ weekly readers for real-life growth advice, playbooks and case studies.

Over 65,000 subscribers

Type your email...

By subscribing, I agree to Substack's [Terms of Use](#), and acknowledge its [Information Collection Notice](#) and [Privacy Policy](#).

Subscribe

Continue reading >

Sign in

GROWTH UNHINGED with Kyle Poyar

Intercom was a first-mover last year with their Fin AI customer support agent and pricing model of \$0.99 per resolution. In Intercom's [words](#), "this ensures that you or pay when Fin does what you care about most; resolving a customer's question."

Zendesk followed with a [similar pricing model](#) based on the number of successful autonomous resolutions by its AI agents. Customers can opt for a flexible pay-as-you-go plan or receive discounts for an upfront commitment.

Each product category has its own flavor of disruptive pricing.

- Legal AI products might charge for a demand package generated by AI or an AI-generated summary.
- Creator AI products might charge for the content that gets produced such as a video generation or amount of video created.
- GTM products might charge for specific tasks completed or workflows executed by the AI.

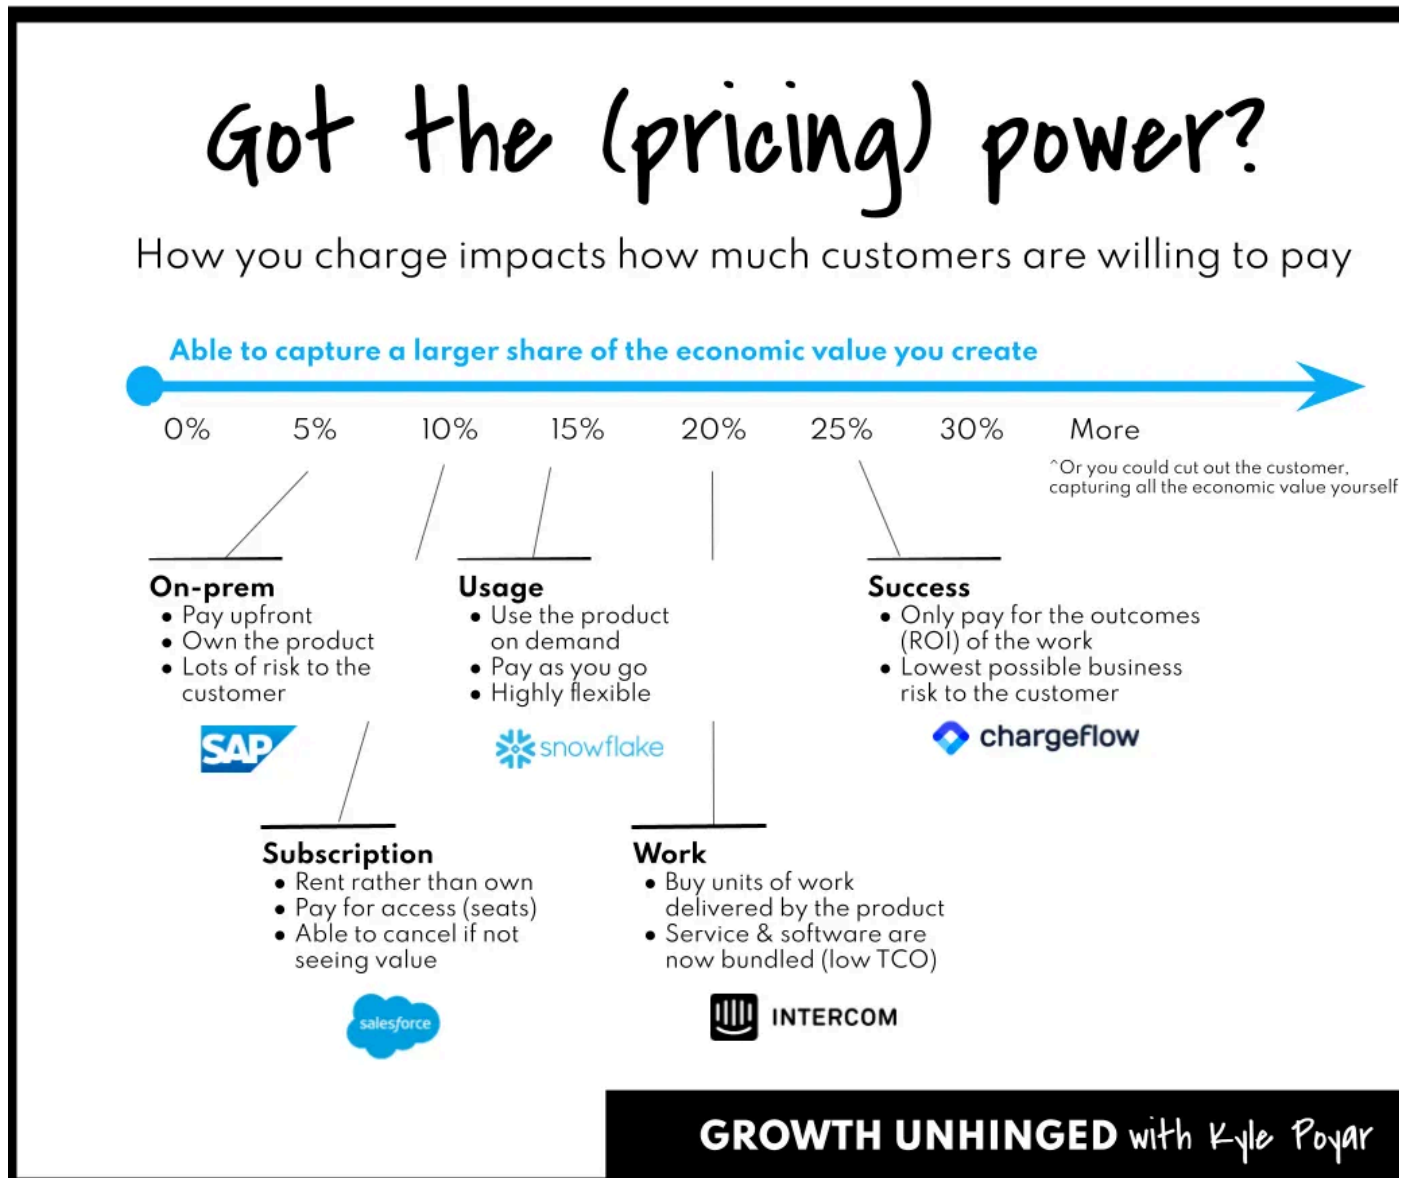
In many cases there's a notion of AI "credits" where more advanced workflows cost more than basic ones. (Some, like workflow task automation provider Bardeen, are [getting creative](#) with how they help customers predict credit consumption.)

Selling work, not necessarily success

In many ways, these disruptive models are a spin on [usage-based](#) or consumption-based pricing where customers pay for a product based on how much they use it. But there are important nuances at play.

Folks are increasingly selling **units of work completed** rather than selling access to the software (seat licenses) or consumption of the software (usage). With the software and the service associated with the software now bundled together, there's the

potential for a lower total cost of ownership (TCO) for the customer along with greater pricing power for the vendor.



Putting this in context, the early iteration of software pricing was the on-premise model. All of the financial risk was on the customer. They paid upfront to own the software — and then were pretty much left on their own until there was a new version to sell.

Shifting to subscriptions lowered the upfront cost — buyers rented rather than own software — and lowered the risk since customers could cancel if they didn't see value.

This made software more accessible, but it didn't necessarily make software cheaper from a lifetime value perspective (I'd argue that this was part of the appeal to investors).

In my experience, subscription SaaS products capture about 10-15% of the estimated economic value they claim to deliver to their customers. Those with success-based billing can capture closer to 20-30% because there's a more direct correlation between the product and the result. Output-based pricing, where customers pay for units of work delivered, falls somewhere in between.

The dark side of success-based pricing

As a customer, I wish I only had to pay for software when it delivered results. But the reality is that true success-based billing won't work for the vast majority of today's products. Most products should charge for work output instead.

The issue is attribution. You want the customer to get a fantastic outcome — and you want them to recognize that your product powered that outcome. As soon as you start charging for success, the customer begins to rethink the results. Did your product really drive the outcome? Or did *they* drive the outcome with a small assist from the product?

If your product relies on people to change their behavior in order to generate ROI, then success-based billing could set you (way) back. Ensure that your product is able to own the service end-to-end and that you're able to align on measurement upfront.

Let's go back to Fin, Intercom's AI agent, which does have a flavor of success-based billing (\$0.99 per AI resolution). You can imagine the potential for attribution fights when the bill comes due.

To Intercom's credit, they have a very thoughtful approach to attribution which they document in their [help center](#). The customer needs to either explicitly confirm that the AI answer is satisfactory (hard resolution) or exit the conversation without returning it within 24 hours (soft resolution). And to give customers even more peace of mind Intercom lets them set usage reminders and hard usage limits to stay in control of their spend.

Resolution definition

A resolution is counted when, following Fin's last answer in a conversation, the customer confirms the answer provided is satisfactory (hard resolution) or exits the conversation without requesting further assistance (soft resolution).

Customers can confirm an answer is satisfactory by:

- Pressing the `That helped` 👍 quick reply.
- Entering an affirmative response such as 'Yes', 'Sure' etc.

Customers can request further assistance by:

- Pressing the `Get more help` quick reply.
- Entering a response that suggests their query is unresolved.
- This includes asking clarifying questions or asking to speak to human support.

A customer has 24 hours to come back to the conversation where Fin provided the answer and select if it helped or not. If after 24 hours the customer did not request further assistance, this will be considered a soft resolution.

Goodbye ARR as we know it?

Shifting to these newer value-based pricing models isn't a simple pricing change you can just announce in a press release.

It's a business model evolution that looks a lot like the shift from on-prem to SaaS in the first place. Things are about to get interesting.

These new AI pricing models might mean greater volatility in both usage and spend. Variable margin profiles across products and customers. Seasonal revenue fluctuations. The potential for project-based, non-recurring use cases.

Put simply, annual recurring revenue (ARR) continues to get dethroned. The new "A" has an extra "R": annual revenue run-rate or ARRR. (For more on this topic, I recommend Dave Kellogg's fantastic presentation about the [impact of AI on SaaS metrics](#).)

Some implications of this shift include:

- We'll need to spend much more time unpacking the **components of revenue**, along with the quality of those revenue line items.
- We'll shift **away from ARR multiple valuations** and instead look at last twelve month (LTM) revenue or, even better, LTM margin dollars.
- We'll pay closer attention to **revenue concentration** as I suspect there will be a far wider variance between the smallest and largest accounts.
- We'll measure newer things like **time to ramp** and **share of wallet** as predictors of future success.

The practical implications

Imagine you're Salesforce making this switch at \$35B in revenue:

- FROM: \$100k+ ACV deals priced per seat with multi-year commits; every seat is the same price regardless of role or usage
- TO: Consumption-based pricing where customers only pay when AI delivers an output, starting at \$2 per conversation

Ignore the strategy side of this. Consider the *practical* realities of making this switch a microcosm for where the industry is headed.

1. The hard work now starts at contract close, it doesn't end there.

This is like shifting from being Planet Fitness to being Barry's Bootcamp. Every day is now a potential revenue event, or a churn event. There's zero room for shelfware or 60% of licenses never logging in.

Your Salesforce rep will likely need to look a lot more like a customer success manager (CSM). And your Salesforce product manager will need to be laser focused on actual product adoption and product-led growth (PLG).

2. Sales comp is about to get, well, messy.

Sales compensation was pretty straightforward in the old model (and quite lucrative for top sellers). Now, what do you do?

You could still pay reps for signing up a big volume commit — but that's pretty backward. It sets up longer deal cycles, over-commitments and future churn. You could look to consumption-based businesses like AWS and Snowflake who pay reps based on actual consumption. Or you could do something more creative like paying on estimated consumption, paying different amounts for booked volume vs. on demand usage, etc.

3. Have fun predicting revenue for Wall Street.

In the old model, forecasting the business was a straightforward spreadsheet exercise with known inputs: pipeline, sales cycles, win rates, and so on. But now you've got to understand (a) actual product adoption, (b) how usage ramps over time in an account, (c) usage seasonality on the customer side, and (d) how improvements to the product allow customers to resolve more conversations. The good news: these changes bring finance leaders much closer to the customer.

4. Look for ways to make usage more palatable to the enterprise.

Traditional enterprise procurement departments are (probably) 0% prepared for an unpredictable bill — even if that bill is better aligned with ROI. As a seller, you'll need to get into the weeds with creative contract structures to make usage more palatable. May I suggest the following:

- **Annual draw-down:** Customers flexibly draw down their usage over 12 months like a gift card. If they use the product faster than expected, they have time to plan and budget before renewal.
- **Grace periods:** After the grace period, customers can either re-up their contract at a higher commit or pay for the one-time flex spend.
- **Roll-over:** Give customers the option of rolling over unused usage credits if the next commit is larger than the last. This helps reduce hoarding behavior.

The industry continues to evolve from selling access to selling value. And this is ultimately a good thing since it opens up new pockets of demand while allowing vendors to better share in customers' success. That's not to say that getting there will be easy.



105 Likes · 7 Restacks

[← Previous](#)[Next →](#)

Discussion about this post

[Comments](#)[Restacks](#)

Write a comment...



Meg Bear Nov 14

❤️ Liked by Kyle Poyar

thanks for writing this - timely as I was midway through writing about this but you covered most points ;-). This is going to create an additional advantage toward disruptors and a lot of complex incumbents.

❤️ LIKE (1) 💬 REPLY



Mark Oehlert Nov 5

❤️ Liked by Kyle Poyar

Great post - critical point for orgs to deeply consider before going in on a solution. Inspired me to write my own post about this so thanks! <https://markoehlert.substack.com/p/agents-and-architecture-and-access>

❤️ LIKE (1) 💬 REPLY



8 more comments...

© 2025 Kyle Poyar · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture